



kaiwanTECH

THE LINUX OPERATING SYSTEM

A VERY BRIEF INTRODUCTION TO IT'S ARCHITECTURE

Important Notice

This courseware is both the product of the author and of freely available opensource materials. Wherever external material has been shown, it's source and ownership have been clearly attributed. We acknowledge all copyrights and trademarks of the respective owners.

The contents of this courseware cannot be copied or reproduced in any form whatsoever without the explicit written consent of the author.

Only the programs - source code and binaries (where applicable) - that form part of this courseware, and that are present on the participant CD, are released under the GNU GPL v2 license and can therefore be used subject to terms of the afore-mentioned license. If you do use any of them, in any manner, you will also be required to clearly attribute their original source (author of this courseware and/or other copyright/trademark holders).

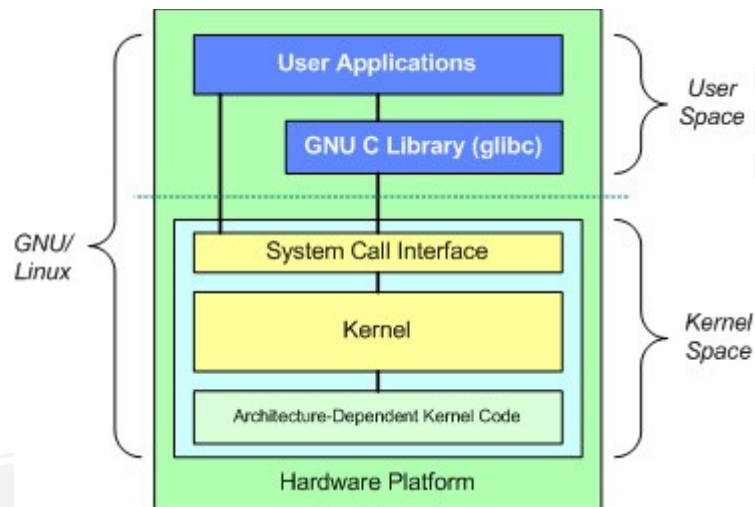
The duration, contents, content matter, programs, etc. contained in this courseware and companion participant CD are subject to change at any point in time without prior notice to individual participants.

Care has been taken in the preparation of this material, but there is no warranty, expressed or implied of any kind and we can assume no responsibility for any errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

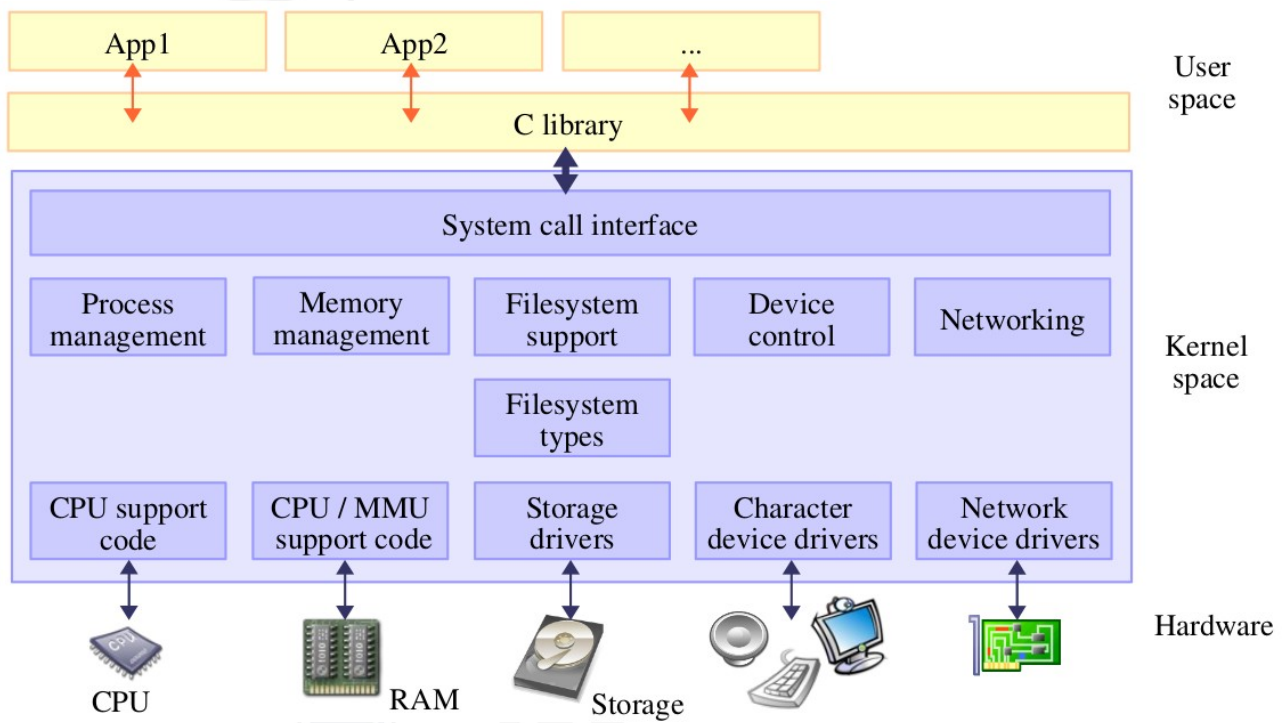
© 2000-2015 Kaiwan N Billimoria
[kaiwanTECH](#), Bangalore, India.

Linux / Unix Architecture

Simplified



More Detailed



<< Above Pic: © Copyright 2006 2004, Michael Opdenacker , © Copyright 2004 2008 Codefidence Ltd. >>

Discussion

- The SCI – System Call Interface – Layer
- CPU Privilege Levels
 - User Mode
 - Kernel (Supervisor) Mode
- Flow of a Process – Birth to Death – between privilege levels

Source**Architecture**

Various layers within Linux, also showing separation between the *userland* and *kernel space*

	User applications	e.g. bash, LibreOffice, Blender, 0 A.D.				
User mode	Low-level system components:	System daemons: <i>systemd</i> , <i>logind</i> , <i>networkd</i> , <i>soundd</i> , ...	Windowing system: <i>display server</i>	Other libraries: <i>GLib</i> , <i>GTK+</i> , <i>Qt</i> , <i>EFL</i> , <i>SDL</i> , <i>SFML</i> , <i>FLTK</i> , <i>GN Ustep</i> , etc.	Graphics: <i>Mesa 3D</i> , <i>AMD Catalyst</i> , ..	
	C standard library	open, exec, sbrk, socket, fopen, calloc, ... (up to 2000 subroutines) <i>glibc</i> aims to be POSIX/SUS-compatible, <i>uClibc</i> targets embedded systems, <i>bionic</i> written for Android, etc.				
		stat, splice, dup, read, open, ioctl, write, mmap, close, exit, etc. (about 380 system calls) The Linux kernel System Call Interface (SCI, aims to be POSIX/SUS-compatible)				
Kernel mode	Linux kernel	Process scheduling subsystem	IPC subsystem	Memory management subsystem	Virtual files subsystem	Network subsystem
		Articles: ALSA, DRI, evdev, LVM, device mapper, Linux Network Scheduler, Netfilter Linux Security Modules: SELinux, TOMOYO, AppArmor, Smack				
Hardware (CPU, main memory, data storage devices, etc.)						

Linux is a **monolithic kernel**. Device drivers and kernel extensions run in **kernel space** (ring 0 in many CPU architectures), with full access to the hardware, although some exceptions run in **user**

[space](#), for example filesystems based on [FUSE](#). The [graphics system](#) most people use with Linux doesn't run in the kernel, in contrast to that found in [Microsoft Windows](#). Unlike standard monolithic kernels, device drivers are easily configured as [modules](#), and loaded or unloaded while running the system. Also unlike standard monolithic kernels, device drivers can be pre-empted under certain conditions. This latter feature was added to handle [hardware interrupts](#) correctly, and to improve support for [symmetric multiprocessing](#).[\[citation needed\]](#) By choice, the Linux kernel has no [Binary Kernel Interface](#).[\[45\]](#)

The hardware is also incorporated into the file hierarchy. Device drivers interface to user applications via an entry in the [/dev](#)[\[46\]](#) and/or [/sys](#) directories. Process information as well is mapped to the file system through the [/proc](#) directory.[\[46\]](#)

Linux supports true [preemptive multitasking](#) (both in [user mode](#) and [kernel mode](#)), [virtual memory](#), [shared libraries](#), [demand loading](#), shared [copy-on-write](#) executables (via [KSM](#)), [memory management](#), the [Internet protocol suite](#), and [threading](#).

...

Some “Golden Rules”

<< To be read later >>

- Modern OS's are Virtual Memory (VM / MMU) – based. Thus:
 - All addresses referred to in user-space are user-mode virtual addresses
 - All addresses referred to in kernel-space are kernel-mode virtual addresses
 - (Very) few exception cases do occur: typically, user-mode device drivers may specify a physical address, kernel-mode memory management is aware of physical addresses and their mapping, some kernel-mode drivers specify a physical (or bus) address.
 - A process can only refer to (lookup) it's own legally-mapped virtual address space: any attempt to look “outside” the process – more correctly – any attempt to lookup an unmapped region result in it being killed by the OS.
 - *Exceptions* to the above:
 - memory-mapping physical memory via `/dev/mem` (need to be root; also see kernel directive `CONFIG_STRICT_DEVMEM`)
 - using the `process_vm_[read|write]v(2)` system calls (kernel ver 3.2 onward [requires kernel support via `CONFIG_CROSS_MEMORY_ATTACH`], glibc 2.15 onward); (need to be root or local and remote processes should have common credentials (IOW, owner)).

[Source](#)

Monolithic Kernel

A monolithic kernel is an operating system architecture where the entire operating system is working in [kernel space](#) and is alone in [supervisor mode](#). The monolithic model differs from other operating system architectures (such as the [microkernel](#) architecture)[\[1\]\[2\]](#) in that it alone defines a high-level virtual interface over computer hardware. A set of primitives or [system calls](#) implement all operating system services such as [process](#) management, [concurrency](#), and [memory management](#). Device drivers can be added to the kernel as [modules](#).

Monolithic architecture examples

[Unix](#) kernels

[BSD](#)
[FreeBSD](#)
[NetBSD](#)
[OpenBSD](#)
[Solaris 1](#) / [SunOS 1.x-4.x](#)

[UNIX System V](#)

[AIX](#)
[HP-UX](#)

[Unix-like](#) kernels

[Linux](#)

[DOS](#)

[DR-DOS](#)
[MS-DOS](#)

____ Microsoft [Windows 9x](#) series ([95](#), [98](#), [Windows 98SE](#), [Me](#))

[OpenVMS](#)

[XTS-400](#)

Microkernel

In [computer science](#), a **microkernel** (also known as μ -kernel) is the near-minimum amount of software that can provide the mechanisms needed to implement an [operating system](#) (OS). These mechanisms include low-level [address space](#) management, [thread](#) management, and [inter-process communication](#) (IPC). If the hardware provides multiple [rings](#) or [CPU modes](#), the microkernel is the only software executing at the most privileged level (generally referred to as [supervisor or kernel mode](#)).

[\[citation needed\]](#) Traditional operating system functions, such as [device drivers](#), [protocol stacks](#) and [file systems](#), are removed from the microkernel to run in [user space](#).[\[citation needed\]](#) In source code size, microkernels tend to be under 10,000 lines of code, as a general rule. [MINIX](#)'s kernel, for

example has fewer than 6,000 lines of code.^[1]

Example : Minix, QNX, VxWorks.

Hybrid

A **hybrid kernel** is a [kernel](#) architecture based on combining aspects of [microkernel](#) and [monolithic kernel](#) architectures used in [computer operating systems](#). The traditional kernel categories are [monolithic kernels](#) and [microkernels](#) (with [nanokernels](#) and [exokernels](#) seen as more extreme versions of microkernels). The category is controversial due to the similarity to monolithic kernel; the term has been dismissed by [Linus Torvalds](#) as simple marketing.^[1]

The idea behind this category is to have a kernel structure similar to a microkernel, but implemented in terms of a monolithic kernel. In contrast to a microkernel, all (or nearly all) operating system services are in [kernel space](#). While there is no performance overhead for message passing and context switching between kernel and user mode, as in [monolithic kernels](#), there are no reliability benefits of having services in [user space](#), as in [microkernels](#).

Implementations

[BeOS](#) kernel

[Haiku](#) kernel

[Syllable](#)

[BSD](#)-based

[DragonFly BSD](#) (first non-[Mach](#) BSD OS to use a hybrid kernel)

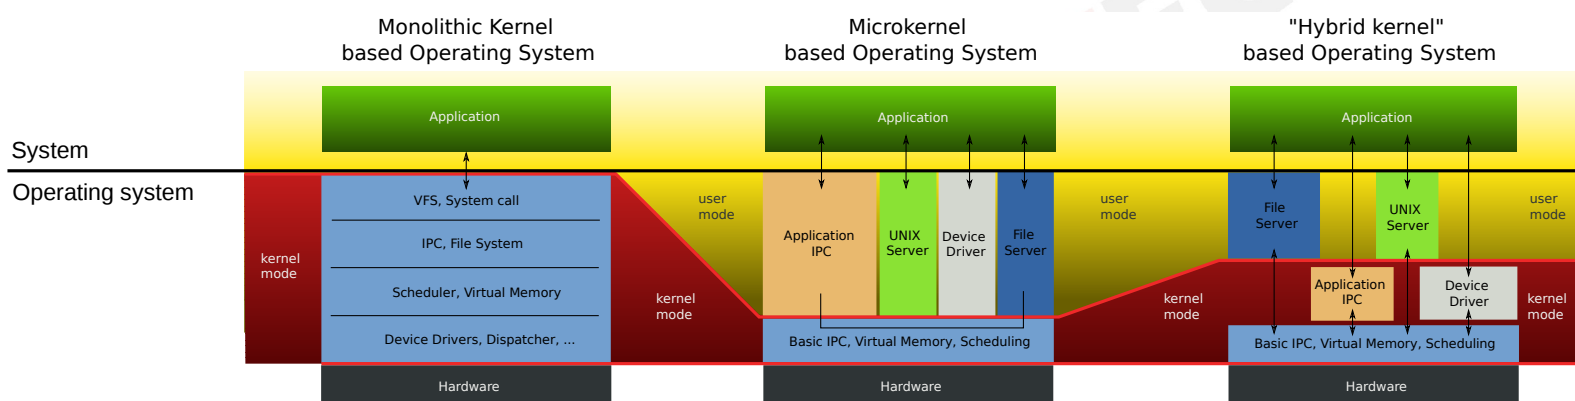
[XNU](#) kernel (core of [Darwin](#), used in [Mac OS X](#) and [iOS](#))

[NetWare](#) kernel^[7]

[Inferno](#) kernel

[NT kernel](#) (used in [Windows NT 3.1](#), [Windows NT 3.5](#), [Windows NT 4.0](#), [Windows 2000](#), [Windows Server 2003](#), [Windows XP](#), [Windows Vista](#), [Windows Server 2008](#), [Windows 7](#), [Windows Server 2008 R2](#), [Windows 8](#), and [Windows Server 2012](#))

[ReactOS](#) kernel



Ref: [Why is Linux called a monolithic kernel?](#)

[P.T.O. -->]

[OPTIONAL / FYI]

An FAQ regarding keeping track of Linux kernel Changes

The Linux kernel is a very fast moving target: things change, quite rapidly at times, new enhancements and features get merged, kernel internal APIs / ABIs get deprecated, etc etc. How can one sanely keep track of all these changes?

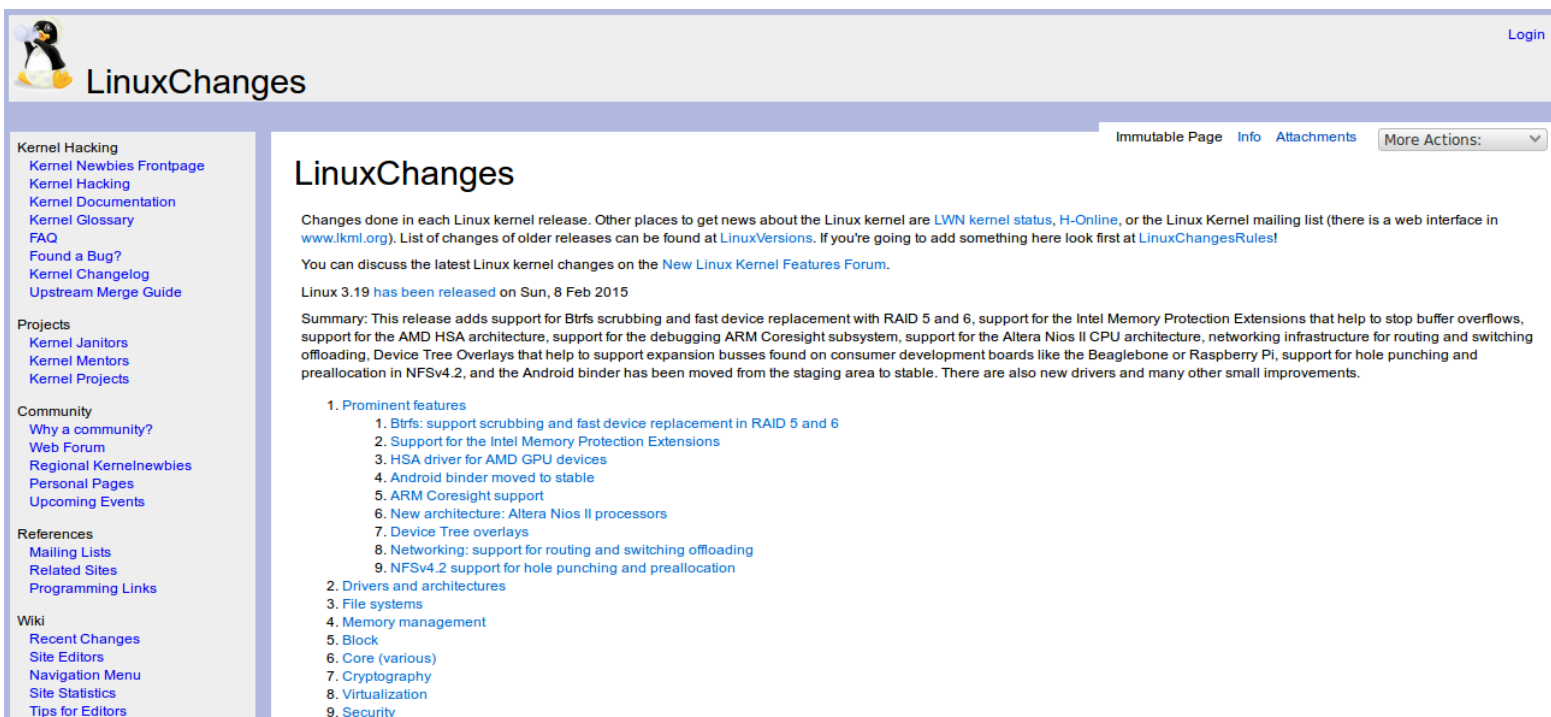
The answer: follow the LKML (Linux Kernel Mailing List).
But “sanely”?
:-)

Read the kernelnewbies “Linux Changes” website !

1. The page

<http://kernelnewbies.org/LinuxChanges>

will have the *latest mainline kernel* changes information:
<< 3.19 at the time of this insertion (March 2015) >>



LinuxChanges

Changes done in each Linux kernel release. Other places to get news about the Linux kernel are [LWN kernel status](#), [H-Online](#), or the Linux Kernel mailing list (there is a web interface in [www.lkml.org](#)). List of changes of older releases can be found at [LinuxVersions](#). If you're going to add something here look first at [LinuxChangesRules!](#)

You can discuss the latest Linux kernel changes on the [New Linux Kernel Features Forum](#).

Linux 3.19 [has been released](#) on Sun, 8 Feb 2015

Summary: This release adds support for Btrfs scrubbing and fast device replacement with RAID 5 and 6, support for the Intel Memory Protection Extensions that help to stop buffer overflows, support for the AMD HSA architecture, support for the debugging ARM Coresight subsystem, support for the Altera Nios II CPU architecture, networking infrastructure for routing and switching offloading, Device Tree Overlays that help to support expansion buses found on consumer development boards like the Beaglebone or Raspberry Pi, support for hole punching and preallocation in NFSv4.2, and the Android binder has been moved from the staging area to stable. There are also new drivers and many other small improvements.

1. Prominent features
 1. Btrfs: support scrubbing and fast device replacement in RAID 5 and 6
 2. Support for the Intel Memory Protection Extensions
 3. HSA driver for AMD GPU devices
 4. Android binder moved to stable
 5. ARM Coresight support
 6. New architecture: Altera Nios II processors
 7. Device Tree overlays
 8. Networking: support for routing and switching offloading
 9. NFSv4.2 support for hole punching and preallocation
2. Drivers and architectures
3. File systems
4. Memory management
5. Block
6. Core (various)
7. Cryptography
8. Virtualization
9. Security

Kernel Hacking

- Kernel Newbies Frontpage
- Kernel Hacking
- Kernel Documentation
- Kernel Glossary
- FAQ
- Found a Bug?
- Kernel Changelog
- Upstream Merge Guide

Projects

- Kernel Janitors
- Kernel Mentors
- Kernel Projects

Community

- Why a community?
- Web Forum
- Regional Kernelnewbies
- Personal Pages
- Upcoming Events

References


- Mailing Lists
- Related Sites
- Programming Links

Wiki

- Recent Changes
- Site Editors
- Navigation Menu
- Site Statistics
- Tips for Editors

Immutable Page Info Attachments More Actions: ▾

2. To see links to all kernel versions, goto
<http://kernelnewbies.org/LinuxVersions>

 **LinuxVersions**

Kernel Hacking
[Kernel Newbies Frontpage](#)
[Kernel Hacking](#)
[Kernel Documentation](#)
[Kernel Glossary](#)
[FAQ](#)
[Found a Bug?](#)
[Kernel Changelog](#)
[Upstream Merge Guide](#)

Projects
[Kernel Janitors](#)
[Kernel Mentors](#)
[Kernel Projects](#)

Community
[Why a community?](#)
[Web Forum](#)
[Regional Kernelnewbies](#)
[Personal Pages](#)
[Upcoming Events](#)

References
[Mailing Lists](#)
[Related Sites](#)
[Programming Links](#)

Wiki
[Recent Changes](#)
[Site Editors](#)
[Navigation Menu](#)
[Site Statistics](#)
[Tips for Editors](#)
[Find Page](#)
[Help Contents](#)
[Hosted by WikiWall](#)

☒ This site ☐ Web

LinuxVersions

This is a list of links to every changelog.


3.x

- [Linux 3.17](#) Released 5 Oct, 2014 (63 days)
- [Linux 3.16](#) Released 3 Aug, 2014 (56 days)
- [Linux 3.15](#) Released 8 June, 2014 (70 days)
- [Linux 3.14](#) Released 30 March, 2014 (70 days)
- [Linux 3.13](#) Released 19 January, 2014 (78 days)
- [Linux 3.12](#) Released 2 November, 2013 (61 days)
- [Linux 3.11](#) Released 2 September, 2013 (64 days)
- [Linux 3.10](#) Released 30 Jun, 2013 (63 days)
- [Linux 3.9](#) Released 28 April, 2013 (69 days)
- [Linux 3.8](#) Released 18 Feb, 2013 70 (days)
- [Linux 3.7](#) Released 10 Dec 2012 (71 days)
- [Linux 3.6](#) Released Sep 30, 2012 (71 days)
- [Linux 3.5](#) Released 21 Jul, 2012 (62 days)
- [Linux 3.4](#) Released 20 May, 2012 (63 days)
- [Linux 3.3](#) Released 18 Mar, 2012 (74 days)
- [Linux 3.2](#) Released 4 Jan, 2012 (72 days)
- [Linux 3.1](#) Released 24 Oct, 2011 (95 days)
- [Linux 3.0](#) Released 21 Jul, 2011 (64 days)

2.6.x

- [Linux 2.6.39](#) Released 18 May, 2011 (65 days)
- [Linux 2.6.38](#) Released 14 March, 2011 (69 days)

Linux Operating System Specialized




The highest quality Training on:

Linux Fundamentals, CLI and Scripting
Linux Systems Programming
Linux Kernel Internals
Linux Device Drivers
Embedded Linux
Linux Debugging Techniques
New! *Linux OS for Technical Managers*

Please do visit our website for details:

<http://kaiwantech.in>



<http://kaiwantech.in>